

porter-neural-networks-regression

April 18, 2024

Porter: Neural Networks Regression By Ratnesh Kumar

Context:

Porter is India's Largest Marketplace for Intra-City Logistics. Leader in the country's \$40 billion intra-city logistics market, Porter strives to improve the lives of 1,50,000+ driver-partners by providing them with consistent earning & independence. Currently, the company has serviced 5+ million customers

Porter works with a wide range of restaurants for delivering their items directly to the people.

Porter has a number of delivery partners available for delivering the food, from various restaurants and wants to get an estimated delivery time that it can provide the customers on the basis of what they are ordering, from where and also the delivery partners.

This dataset has the required data to train a regression model that will do the delivery time estimation, based on all those features

Data Dictionary

Each row in this file corresponds to one unique delivery. Each column corresponds to a feature as explained below.

- market_id : integer id for the market where the restaurant lies
- created_at : the timestamp at which the order was placed
- actual_delivery_time : the timestamp when the order was delivered
- store_primary_category : category for the restaurant
- order_protocol : integer code value for order protocol(how the order was placed ie: through porter, call to restaurant, pre booked, third part etc)
- total_items
- subtotal : final price of the order
- num_distinct_items : the number of distinct items in the order
- min_item_price : price of the cheapest item in the order
- max_item_price : price of the costliest item in order
- total_onshift_partners : number of delivery partners on duty at the time order was placed
- total_busy_partners : number of delivery partners attending to other tasks
- total_outstanding_orders : total number of orders to be fulfilled at the moment

```
[12]: import pandas as pd
df=pd.read_csv('/content/sample_data/Porterdataset.csv')
df.head()
```

| | market_id | created_at | actual_delivery_time | |
|---|-----------|---------------------|----------------------|--|
| 0 | 1.0 | 2015-02-06 22:24:17 | 2015-02-06 23:27:16 | |
| 1 | 2.0 | 2015-02-10 21:49:25 | 2015-02-10 22:56:29 | |
| 2 | 3.0 | 2015-01-22 20:39:28 | 2015-01-22 21:09:09 | |
| 3 | 3.0 | 2015-02-03 21:21:45 | 2015-02-03 22:13:00 | |
| 4 | 3.0 | 2015-02-15 02:40:36 | 2015-02-15 03:20:26 | |

| | store_id | store_primary_category | order_protocol | |
|---|----------------------------------|------------------------|----------------|--|
| 0 | df263d996281d984952c07998dc54358 | american | 1.0 | |
| 1 | f0ade77b43923b38237db569b016ba25 | mexican | 2.0 | |
| 2 | f0ade77b43923b38237db569b016ba25 | NaN | 1.0 | |
| 3 | f0ade77b43923b38237db569b016ba25 | NaN | 1.0 | |
| 4 | f0ade77b43923b38237db569b016ba25 | NaN | 1.0 | |

| | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | |
|---|-------------|----------|--------------------|----------------|----------------|--|
| 0 | 4 | 3441 | 4 | 557 | 1239 | |
| 1 | 1 | 1900 | 1 | 1400 | 1400 | |
| 2 | 1 | 1900 | 1 | 1900 | 1900 | |
| 3 | 6 | 6900 | 5 | 600 | 1800 | |
| 4 | 3 | 3900 | 3 | 1100 | 1600 | |

| | total_onshift_partners | total_busy_partners | total_outstanding_orders | |
|---|------------------------|---------------------|--------------------------|--|
| 0 | 33.0 | 14.0 | 21.0 | |
| 1 | 1.0 | 2.0 | 2.0 | |
| 2 | 1.0 | 0.0 | 0.0 | |
| 3 | 1.0 | 1.0 | 2.0 | |
| 4 | 6.0 | 6.0 | 9.0 | |

Defining problem statement, importing the data and data structure analysis (10 points)

PROBLEM STATEMENT: Porter has a number of delivery partners available for delivering the food, from various restaurants and wants to get an estimated delivery time that it can provide the customers on the basis of what they are ordering, from where and also the delivery partners.

This dataset has the required data to train a regression model that will do the delivery time estimation, based on all those features

```
[13]: df.info()
```

| # | Column | Non-Null Count | Dtype |
|-----|--------|----------------|-------|
| --- | --- | ----- | ----- |

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   market_id        197428 non-null   float64
 1   created_at       197428 non-null   datetime64[ns]
 2   actual_delivery_time  197428 non-null   datetime64[ns]
 3   store_id          197428 non-null   object  
 4   store_primary_category  197428 non-null   object  
 5   order_protocol    197428 non-null   float64
 6   total_items       197428 non-null   int64  
 7   subtotal          197428 non-null   int64  
 8   num_distinct_items 197428 non-null   int64  
 9   min_item_price    197428 non-null   float64
 10  max_item_price    197428 non-null   float64
 11  total_onshift_partners 197428 non-null   float64
 12  total_busy_partners 197428 non-null   float64
 13  total_outstanding_orders 197428 non-null   float64
```

```

0   market_id           196441 non-null  float64
1   created_at          197428 non-null  object
2   actual_delivery_time 197421 non-null  object
3   store_id             197428 non-null  object
4   store_primary_category 192668 non-null  object
5   order_protocol        196433 non-null  float64
6   total_items           197428 non-null  int64
7   subtotal              197428 non-null  int64
8   num_distinct_items    197428 non-null  int64
9   min_item_price         197428 non-null  int64
10  max_item_price        197428 non-null  int64
11  total_onshift_partners 181166 non-null  float64
12  total_busy_partners    181166 non-null  float64
13  total_outstanding_orders 181166 non-null  float64
dtypes: float64(5), int64(5), object(4)
memory usage: 21.1+ MB

```

[14]: df.isna().sum()

```

[14]: market_id           987
       created_at          0
       actual_delivery_time 7
       store_id             0
       store_primary_category 4760
       order_protocol        995
       total_items           0
       subtotal              0
       num_distinct_items    0
       min_item_price         0
       max_item_price         0
       total_onshift_partners 16262
       total_busy_partners    16262
       total_outstanding_orders 16262
dtype: int64

```

```

[15]: print('NAN in total_onshift_partners :
      ↵',(len(df['total_onshift_partners'])-df['total_onshift_partners'].isna() .
      ↵sum())/len(df['total_onshift_partners'])*100)
df.drop(['total_onshift_partners'],inplace=True,axis=1)
print('NAN in total_busy_partners :
      ↵',(len(df['total_busy_partners'])-df['total_busy_partners'].isna().sum())/
      ↵len(df['total_busy_partners'])*100)
df.drop(['total_busy_partners'],inplace=True,axis=1)
print('NAN in total_outstanding_orders :
      ↵',(len(df['total_outstanding_orders'])-df['total_outstanding_orders'].isna() .
      ↵sum())/len(df['total_outstanding_orders'])*100)
df.drop(['total_outstanding_orders'],inplace=True,axis=1)

```

```
df.head()
```

```
NAN in total_onshift_partners : 91.7630731203274  
NAN in total_busy_partners : 91.7630731203274  
NAN in total_outstanding_orders : 91.7630731203274
```

```
[15]:    market_id      created_at actual_delivery_time \
0        1.0  2015-02-06 22:24:17  2015-02-06 23:27:16
1        2.0  2015-02-10 21:49:25  2015-02-10 22:56:29
2        3.0  2015-01-22 20:39:28  2015-01-22 21:09:09
3        3.0  2015-02-03 21:21:45  2015-02-03 22:13:00
4        3.0  2015-02-15 02:40:36  2015-02-15 03:20:26

                           store_id store_primary_category order_protocol \
0  df263d996281d984952c07998dc54358                  american      1.0
1  f0ade77b43923b38237db569b016ba25                mexican      2.0
2  f0ade77b43923b38237db569b016ba25                  NaN      1.0
3  f0ade77b43923b38237db569b016ba25                  NaN      1.0
4  f0ade77b43923b38237db569b016ba25                  NaN      1.0

   total_items  subtotal  num_distinct_items  min_item_price  max_item_price
0          4       3441                  4           557            1239
1          1       1900                  1           1400            1400
2          1       1900                  1           1900            1900
3          6       6900                  5           600            1800
4          3       3900                  3           1100            1600
```

** Data preprocessing and feature engineering (30 points)

- Data cleaning
- Null value handling
- Creating the target column (time taken for delivery) from order timestamp and delivery timestamp
- Getting hour and day of the week
- Encoding categorical column

**

```
[16]: df['store_primary_category'].value_counts()
df['store_primary_category'].fillna('Other', inplace=True)
df['market_id'].value_counts()
df['market_id'].fillna('0.0', inplace=True)
df['order_protocol'].value_counts()
df['order_protocol'].fillna('0.0', inplace=True)
df.dropna(subset=['actual_delivery_time'], inplace=True)
df.isna().sum()
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 197421 entries, 0 to 197427
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   market_id        197421 non-null   object 
 1   created_at       197421 non-null   object 
 2   actual_delivery_time 197421 non-null   object 
 3   store_id         197421 non-null   object 
 4   store_primary_category 197421 non-null   object 
 5   order_protocol   197421 non-null   object 
 6   total_items      197421 non-null   int64  
 7   subtotal         197421 non-null   int64  
 8   num_distinct_items 197421 non-null   int64  
 9   min_item_price   197421 non-null   int64  
 10  max_item_price   197421 non-null   int64  
dtypes: int64(5), object(6)
memory usage: 18.1+ MB

```

```
[17]: #Categorical columns
cat_col = [col for col in df.columns if df[col].dtype == 'object']
print('Categorical columns :',cat_col)
# Numerical columns
num_col = [col for col in df.columns if df[col].dtype != 'object']
print('Numerical columns :',num_col)
df[cat_col].nunique()
```

```
Categorical columns : ['market_id', 'created_at', 'actual_delivery_time',
'store_id', 'store_primary_category', 'order_protocol']
Numerical columns : ['total_items', 'subtotal', 'num_distinct_items',
'min_item_price', 'max_item_price']
```

```
[17]: market_id          7
      created_at        180981
      actual_delivery_time 178110
      store_id          6743
      store_primary_category 75
      order_protocol     8
      dtype: int64
```

```
[18]: df['actual_delivery_time']=df['actual_delivery_time'].astype('datetime64[ns]')
df['created_at']=df['created_at'].astype('datetime64[ns]')
df['Time_taken_for_delivery']= (df['actual_delivery_time']-df['created_at'])/pd.
    .Timedelta('60s')
df['hour']=df['created_at'].dt.hour
df['day']=df['created_at'].dt.dayofweek
```

Dropping the Column that are no longer required

```
[19]: df.drop(['created_at', 'actual_delivery_time', 'store_id'], axis=1, inplace=True)
```

```
[20]: df['store_primary_category']=df['store_primary_category'].astype('category').  
      ↪cat.codes  
  
df.head()
```

```
[20]:   market_id  store_primary_category  order_protocol  total_items  subtotal  \  
0          1.0                      5            1.0           4       3441  
1          2.0                     48            2.0           1       1900  
2          3.0                      0            1.0           1       1900  
3          3.0                      0            1.0           6      6900  
4          3.0                      0            1.0           3      3900  
  
      num_distinct_items  min_item_price  max_item_price  \  
0                  4             557          1239  
1                  1            1400          1400  
2                  1            1900          1900  
3                  5             600          1800  
4                  3            1100          1600  
  
      Time_taken_for_delivery  hour  day  
0            62.983333     22    4  
1            67.066667     21    1  
2            29.683333     20    3  
3            51.250000     21    1  
4            39.833333      2    6
```


Data visualization and cleaning (20 points)

- Visualization for various features
- Check for outliers
- Remove outliers
- Compare plots and results

```
[21]: import seaborn as sns  
import matplotlib.pyplot as plt  
sns.heatmap(df.corr())
```


```
<ipython-input-21-fbcf84b297f6>:3: FutureWarning: The default value of  
numeric_only in DataFrame.corr is deprecated. In a future version, it will  
default to False. Select only valid columns or specify the value of numeric_only  
to silence this warning.  
  sns.heatmap(df.corr())
```

```
[21]: <Axes: >
```




```
[22]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(10,2))
plt.subplot(1,2,1)
sns.barplot(y='total_items',x='market_id',data=df,estimator='sum')
plt.subplot(1,2,2)
sns.barplot(y='total_items',x='store_primary_category',data=df,estimator='sum')
```

```
[22]: <Axes: xlabel='store_primary_category', ylabel='total_items'>
```




```
[23]: plt.figure(figsize=(5,2))
sns.barplot(y='total_items',x='order_protocol',data=df,estimator='sum')
```

```
[23]: <Axes: xlabel='order_protocol', ylabel='total_items'>
```




```
[24]: plt.figure(figsize=(10,2))
plt.subplot(121)
sns.countplot(x=df['day'])
plt.subplot(122)
sns.countplot(x=df['hour'])
```

```
[24]: <Axes: xlabel='hour', ylabel='count'>
```



```
[25]: sns.scatterplot(x='hour',y='Time_taken_for_delivery',data=df)
```

```
[25]: <Axes: xlabel='hour', ylabel='Time_taken_for_delivery'>
```




[25]:

No Collinearity

Detecting Outliers

```
[27]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(2,2))
sns.boxplot(y='Time_taken_for_delivery',data=df)
plt.xticks(rotation=90);
plt.show()
```




Removing Outliers

```
[28]: import numpy as np  
print((df.loc[df['Time_taken_for_delivery'] >400].shape[0] / df.shape[0]) * 100)  
df.drop(index=df.loc[df['Time_taken_for_delivery'] >400].index[0],inplace=True)
```

0.009624102805679234

```
[29]: import seaborn as sns  
import matplotlib.pyplot as plt  
plt.figure(figsize=(2,2))  
sns.boxplot(y='Time_taken_for_delivery',data=df)  
plt.xticks(rotation=90);  
plt.show()
```



Model training with random forest (10 points)

- Data splitting
- Random forest regression

```
[33]: from sklearn.model_selection import train_test_split  
from sklearn.ensemble import RandomForestRegressor  
y=df['Time_taken_for_delivery']  
x=df.drop(['Time_taken_for_delivery'],axis=1)  
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.  
    ↪2,random_state=42)  
regressor=RandomForestRegressor()  
regressor.fit(X_train,y_train)
```

```
[33]: RandomForestRegressor()
```

```
[40]: from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score  
prediction=regressor.predict(X_test)
```

```

mse=mean_squared_error(y_test,prediction)
rmse=mse**.5
print("mse : ",mse)
print("rmse : ",rmse)
mae=mean_absolute_error(y_test,prediction)
print("mase : ",mae)
mape=np.mean(np.abs((y_test - prediction)/y_test))*100
print("mape : ",mape)

r2_score(y_test,prediction)

```

```

mse : 1298.7022878370515
rmse : 36.037512231521355
mase : 12.831212676817074
mape : 29.537370372638787

```

[40]: 0.018423261039891337

Regression with neural networks (30 points)

- Data scaling
- Defining NN architecture
- Trying different combinations and hyperparameters
- Model training
- Comparing results with random forest

[41]: #Scaling the data to feed before neural network

```

from sklearn import preprocessing
scaler=preprocessing.MinMaxScaler()
x_scaled=scaler.fit_transform(x)
X_train,X_test,y_train,y_test=train_test_split(x_scaled,y,test_size=0.
                                               ↴2,random_state=42)

```

We will build a simple neural network to train our regression model it is a sequential model with two layers,

we have kept the number of nodes in the first layers equal to the number of input columns, and for the subsequent layers 32, 32, which can we changed or experimented with

the activation for the layers is kept as relu because it is a great non linear activation function that works for most cases, we could have used leaky relu if we see gradient vanishing.

the last layer has one node because it will give the single result that is our delivery time and the activation function for that should be linear

[74]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

```

model=Sequential()
model.add(Dense(11,kernel_initializer='normal',activation='relu'))
model.add(Dense(32,activation='relu'))
model.add(Dense(32,activation='relu'))
model.add(Dense(1,activation='linear'))
model.compile(loss='mse',optimizer='Adam',metrics=['mse','mae'])
history=model.
    ↪fit(X_train,y_train,epochs=10,batch_size=512,verbose=1,validation_split=0.2)

```

Epoch 1/10
247/247 [=====] - 3s 6ms/step - loss: 1250.7706 - mse:
1250.7706 - mae: 27.2805 - val_loss: 1368.2872 - val_mse: 1368.2872 - val_mae:
14.8899
Epoch 2/10
247/247 [=====] - 1s 4ms/step - loss: 422.8232 - mse:
422.8232 - mae: 14.2876 - val_loss: 1322.7933 - val_mse: 1322.7933 - val_mae:
13.9953
Epoch 3/10
247/247 [=====] - 1s 4ms/step - loss: 390.1947 - mse:
390.1947 - mae: 13.6241 - val_loss: 1297.1119 - val_mse: 1297.1119 - val_mae:
13.4820
Epoch 4/10
247/247 [=====] - 1s 4ms/step - loss: 373.6034 - mse:
373.6034 - mae: 13.2969 - val_loss: 1285.8724 - val_mse: 1285.8724 - val_mae:
13.3453
Epoch 5/10
247/247 [=====] - 1s 4ms/step - loss: 368.1090 - mse:
368.1090 - mae: 13.2022 - val_loss: 1283.5146 - val_mse: 1283.5146 - val_mae:
13.3899
Epoch 6/10
247/247 [=====] - 1s 3ms/step - loss: 366.8733 - mse:
366.8733 - mae: 13.1980 - val_loss: 1282.6769 - val_mse: 1282.6769 - val_mae:
13.2896
Epoch 7/10
247/247 [=====] - 1s 3ms/step - loss: 366.4678 - mse:
366.4678 - mae: 13.1877 - val_loss: 1283.0757 - val_mse: 1283.0757 - val_mae:
13.4240
Epoch 8/10
247/247 [=====] - 1s 3ms/step - loss: 366.3302 - mse:
366.3302 - mae: 13.1955 - val_loss: 1283.0736 - val_mse: 1283.0736 - val_mae:
13.1529
Epoch 9/10
247/247 [=====] - 1s 3ms/step - loss: 366.1603 - mse:
366.1603 - mae: 13.1877 - val_loss: 1282.5369 - val_mse: 1282.5369 - val_mae:
13.2340
Epoch 10/10


```
247/247 [=====] - 1s 3ms/step - loss: 366.2173 - mse:  
366.2173 - mae: 13.1884 - val_loss: 1283.1187 - val_mse: 1283.1187 - val_mae:  
13.1472
```

```
[75]: model.summary()  
from tensorflow.keras.utils import plot_model  
plot_model(model)
```


```
Model: "sequential_26"
```

| Layer (type) | Output Shape | Param # |
|-------------------------------------|--------------|---------|
| <hr/> | | |
| dense_120 (Dense) | (None, 11) | 121 |
| dense_121 (Dense) | (None, 32) | 384 |
| dense_122 (Dense) | (None, 32) | 1056 |
| dense_123 (Dense) | (None, 1) | 33 |
| <hr/> | | |
| Total params: 1594 (6.23 KB) | | |
| Trainable params: 1594 (6.23 KB) | | |
| Non-trainable params: 0 (0.00 Byte) | | |

```
[75]:
```



```
[78]: plt.plot(history.history['loss'])
      plt.plot(history.history['val_loss'])
      plt.xlabel("Epochs")
      plt.ylabel('loss')
      plt.legend(['loss', 'val_loss'])
      plt.show()
```



```
[84]: print('r2_score:', r2_score(y_test, model.predict(X_test)))
mse = mean_squared_error(y_test, model.predict(X_test))
rmse = mse**.5
print("mse : ",mse)
print("rmse : ",rmse)
print("errors for neural net")
mae = mean_absolute_error(y_test, model.predict(X_test))
print("mae : ",mae)
```

1234/1234 [=====] - 4s 3ms/step
r2_score: 0.01630627602202872
1234/1234 [=====] - 2s 2ms/step
mse : 1301.5032234917896
rmse : 36.07635269108824
errors for neural net
1234/1234 [=====] - 2s 1ms/step
mae : 13.135253871507697

```
[85]: from sklearn.metrics import mean_absolute_percentage_error
mean_absolute_percentage_error(y_test, model.predict(X_test))
```

1234/1234 [=====] - 3s 2ms/step

[85]: 0.29956791711312275

[]: